

# Project Management — Koorde vs Chord Benchmark

## 1. Project Overview & Inspiration

This project evaluates the performance of distributed hash tables (DHTs) in a realistic cloud environment. It was inspired by:

1. **The Koorde Paper (Kaashoek & Karger, 2003):** Proposed a degree-optimal DHT with  $\mathcal{O}(\log N / \log \log N)$  diameter.
2. **Existing Benchmarks:** Online comparisons of Chord vs. Kademia, (notably <https://github.com/macvincent/slbdch>), which we aimed to extend by validating Koorde's theoretical properties.

**Goal:** Move beyond simple algorithmic simulation to a full systems engineering evaluation, measuring **latency**, **cache hit rate**, and **throughput** in a containerized, orchestrated environment.

## 2. Team Roles & Responsibilities

To manage the complexity of a distributed systems project, we divided responsibilities based on the engineering stack layers.

Role	Member	Key Responsibilities
Protocol Engineer	Yuzheng Shi	<b>Core Logic &amp; Correctness:</b> <ul style="list-style-type: none"><li>Implemented Koorde routing (de Bruijn graph) &amp; Chord finger tables.</li><li>Validated routing correctness and stabilization logic.</li><li><b>Experiments:</b> Conducted Latency and Cache Hit Rate (Churn) experiments.</li></ul>
Infrastructure Lead	Anran Lyu	<b>Deployment &amp; Scale:</b> <ul style="list-style-type: none"><li>Containerized the application (Docker) and designed K8s manifests.</li><li>Managed AWS EKS infrastructure and LocalStack testing.</li><li><b>Experiments:</b> Led Throughput/Saturation testing and distributed Locust load generation.</li></ul>

## 3. Project Timeline & Milestones

The project followed an iterative engineering lifecycle over the course of the term.

Phase	Milestone	Description	Status
Phase 1	Core Implementation	Implemented <code>DHTNode</code> interface, Koorde logic ( <code>logicnode</code> ), and Chord logic ( <code>chord</code> ). Verified routing locally.	<span style="color: green;">✓</span> Completed
Phase 2	Local Integration	Built Docker images ( <code>node.Dockerfile</code> ) and created <code>docker-compose</code> / LocalStack setups to test multi-node interaction.	<span style="color: green;">✓</span> Completed
Phase 3	Cloud Deployment	Migrated to AWS EKS. Solved IAM and quota issues. Established a stable 8-32 node cluster environment.	<span style="color: green;">✓</span> Completed
Phase 4	Benchmarking	Executed 3 major experiments: Latency Scaling, Churn Resilience, and Throughput Saturation.	<span style="color: green;">✓</span> Completed
Phase 5	Analysis & Reporting	Aggregated CSV results, generated Python plots, and produced the final <code>experience_report.md</code> .	<span style="color: green;">✓</span> Completed

## 4. Architecture & Implementation Breakdown

We structured the repository to separate concerns, allowing parallel development.

### A. Routing Layer (The "Brain")

Owned by Protocol Engineer

- **Koorde:** Implemented in `internal/node/logicnode` and `internal/node/routingtable`. Challenges included managing "imaginary" nodes and de Bruijn graph transitions.
- **Chord:** Implemented in `internal/node/chord`. Focused on standard stabilization and finger table maintenance.

### B. Infrastructure Layer (The "Body")

- **Containerization:** docker/ contains optimized builds for the Node and Client.
- **Orchestration:** deploy/eks contains the Kubernetes manifests (Deployments, Services, ConfigMaps) required to run the cluster in AWS.
- **Automation:** scripts/ contains PowerShell automation (e.g., test-koorde-scaling.ps1 ) to orchestrate complex test scenarios.

## C. Shared Components

*Jointly Developed*

- **Web Cache:** internal/node/cache implements the LRU cache and hotspot detection.
- **RPC Layer:** gRPC definitions in proto/ ensure consistent communication between nodes.

## 5. Risk Management & Adaptations

We encountered significant hurdles that required agile adaptation of our project scope.

Risk / Problem	Impact	Mitigation / Adaptation
<b>AWS Learner Lab Quotas</b>	Hard limit of ~9 EC2 instances prevented planned 100-node physical cluster.	<b>Scope Adjustment:</b> Refocused on detailed analysis of small-to-medium clusters (8-32 nodes). Proposed logical scaling and simulation for asymptotic validation.
<b>Sandbox IAM Restrictions</b>	AWS Student Sandbox lacked IAM roles for EKS, blocking initial deployment.	<b>Platform Migration:</b> Reverted to AWS Learner Lab and optimized resource usage (packing multiple pods per node) to maximize cluster size within limits.
<b>Load Generator Crashing</b>	Local machine CPU exhausted when simulating >1000 users with Locust.	<b>Distributed Testing:</b> Architected a distributed Locust setup on AWS (1 Master + N Workers) to generate stable load up to 4000 users.

## 6. Key Artifacts & Evidence

The following artifacts in the repository demonstrate the completion of our goals:

- **Source Code:**
  - Koorde Logic: internal/node/logicnode/
  - Chord Logic: internal/node/chord/
- **Infrastructure:**
  - K8s Manifests: deploy/eks/
  - Dockerfiles: docker/
- **Data & Analysis:**
  - Benchmark Report: BENCHMARK\_REPORT.md
  - Raw Data: test/results/32nodes/\*.csv
  - Visualization Script: test/generate\_plots.py

## 7. Final Reflection

What began as a theoretical implementation exercise evolved into a comprehensive distributed systems engineering project. We successfully:

- Built a functional, multi-protocol distributed cache.
- Navigated real-world cloud constraints (quotas, IAM).
- Produced empirical data comparing theoretical promises (Koorde) vs. practical reality (Chord).

While we could not demonstrate Koorde's asymptotic superiority due to scale limits, we successfully proved that **Chord is more stable and performant** for small-to-medium clusters, a valuable negative result for system architects.